

Lecture 10: Graph Search

*Harvard SEAS - Fall 2024**2024-10-03*

1 Announcements

- Sender-Receiver exercise on Tuesday.

2 Graph Algorithms

Recommended Reading:

- Roughgarden II Sec 7.0–7.3, 8.0–8.1.1
- CLRS Appendix B.4

Motivating Problem: Google Maps.

Given a road network, a starting point, and a destination, what is the shortest way to get from the starting point s to the destination t ?

Q: How to model a road network?

A:

Definition 2.1.

Q: What possibilities doesn't this model capture and how might we augment it?

Unless we state otherwise, assume *graph* means a **simple, unweighted, undirected** graph, and a *digraph* means a **simple, unweighted, directed** graph.

Although we've started with this motivation of route-finding on road networks, in the coming lectures we'll see that graphs are useful for modelling a vast range of different kinds of relationships, e.g. social networks, the world wide web, kidney donor compatibilities, scheduling conflicts, etc. 50

Representing Graphs.

- Adjacency list representation: For every vertex v , given
 - $\text{Nbr}_{\text{out}}[v] =$
 - $\text{deg}_{\text{out}}(v) =$
- Using Word-RAM with word length $w = O(\log n)$, so vertex name fits in a single word.

3 Shortest Walks

Abstracting a simplified version of the route-finding problem above, we wish to design an algorithm for the following computational problem:

Input	: A digraph $G = (V, E)$ and two vertices $s, t \in V$
Output	: A <i>shortest walk</i> from s to t in G , if any walk from s to t exists

Computational Problem ShortestWalk

Let us define precisely what we mean by a *shortest walk*.

Definition 3.1. Let $G = (V, E)$ be a directed graph, and $s, t \in V$.

- A *walk* w from s to t in G is
- A walk in which all vertices are distinct is also called a *path*.
- The *length* of a walk w is $\text{length}(w) =$
- The *distance* from s to t in G is $\text{dist}_G(s, t) =$
- A *shortest walk* from s to t in G is a walk w from s to t with $\text{length}(w) = \text{dist}_G(s, t)$

Q: What algorithm for ShortestWalk is immediate from the definition?

A: But when can we stop this algorithm to conclude that there is no walk? The following lemma allows us to stop at walks of length $n - 1$.

Lemma 3.2. *If w is a shortest walk from s to t , then all of the vertices that occur on w are distinct (i.e. w is a path).*

Proof.

□

Because of this lemma, the ShortestWalk problem is usually referred to as the *ShortestPath* problem.

Q: With this lemma, what is the runtime of exhaustive search?

A:

4 Breadth-First Search

“I don’t know where I’m going, but I’m on my way.” — Carl Sagan

We can get a faster algorithm using *breadth-first search (BFS)*.

For simplicity, we’ll start by presenting algorithms to only compute the *length* of the shortest path from s to t , rather than actually find the path. On the other hand, our algorithm will actually compute the distance from s to *all* vertices in the graph, not only t . Let’s capture these two modifications in the following definition:

Input	: A digraph $G = (V, E)$ and a source vertex $s \in V$
Output	: The array dist_s where for every $t \in V$, $\text{dist}_s[t] = \text{dist}_G(s, t)$

Computational Problem SingleSourceDistances

With this, here is our first version of BFS.

```
1 BFS( $G, s$ )
   Input           : A digraph  $G = (V, E)$  and a source vertex  $s \in V$ 
   Output          : The array  $\text{dist}_s[\cdot] = \text{dist}_G(s, \cdot)$ 
2 Initialize  $\text{dist}_s[t] = \infty$  for all  $t \in V$ .;
3  $S = F = \{s\}$ ;
4  $\text{dist}_s[s] = 0$ ;
5 foreach  $d = 1, \dots, n - 1$  do
6   | Let  $F = \{v \in V : v \notin S, \exists u \in F \text{ s.t. } (u, v) \in E\}$ ;
7   | For every  $v \in F$ ,  $\text{dist}_s[v] = d$ ;
8   |  $S = S \cup F$ ;
9 return  $\text{dist}_s$ 
10
```

Algorithm 1: BFS for SingleSourceDistances

Example:

Q: What is happening at every iteration of the loop? We have a set of S which is the set of vertices that have been visited previously, and F the *frontier*, which is the set of vertices that were visited for the first time in the previous iteration. At each iteration, we update F by taking all vertices not previously visited that can be reached from the previous frontier by one additional edge. Then we add all of these new frontier vertices into S .

Q: How do we prove correctness?

Q: What is the runtime of the algorithm, in terms of the number of vertices n and the number of edges m ?

Q: How can we implement BFS faster?

Putting all the above together, we obtain:

Theorem 4.1. *BFS(G) correctly solves SingleSourceDistances and can be implemented in time $O(n + m)$, where n is the number of vertices in G and m is the number of edges in G .*

Implementation details:

- In practice, F is stored as a queue and updates are done one vertex u at a time rather than as a ‘batch’.
- The bitvector S is redundant given that we are maintaining \mathbf{dist}_s .

5 Finding the Paths

Q: How to actually find shortest *path* from s to t , not just the distances?

Input	: A digraph $G = (V, E)$ and a source vertex $s \in V$
Queries	: For any query vertex t , return a shortest path from s to t (if one exists).

Data-Structure Problem SingleSourceShortestPaths

Theorem 5.1. *There is a solution that solves the SingleSourceShortestPaths problem on digraphs with n vertices and m edges with Preprocessing time $O(n + m)$ and the time to answer a query t is $O(\text{dist}_G(s, t))$.*

It is natural and very useful (e.g. in Google maps) to have data structures for Shortest Paths on *dynamic graphs*. There is a rich collection of methods for dynamic graph data structures, which are beyond the scope of this course.

Another extension of BFS is to handle *weighted graphs*. This is Dijkstra's algorithm, and is covered in CS 1240.

6 (Optional) Other Forms of Graph Search

Another very useful form of graph search that you may have seen is *depth-first search* (DFS). We won't cover it in CS120, but DFS and some of its applications are covered in CS124.

We do, however, briefly mention a randomized form of graph search, namely *random walks*, and use it to solve the *decision* problem of STConnectivity on undirected graphs.

Input	: A graph $G = (V, E)$ and vertices $s, t \in V$
Output	: YES if there is a walk from s to t in G , and NO otherwise

Computational Problem UndirectedSTconnectivity

```

1 RandomWalk( $G, s, \ell$ )
   Input      : A digraph  $G = (V, E)$ , a vertices  $s, t \in V$ , and a walk-length  $\ell$ 
   Output    : YES or NO
2  $v = s$ ;
3 foreach  $i = 1, \dots, \ell$  do
4   | if  $v = t$  then return YES;
5   |  $j = \text{random}(\text{deg}_{\text{out}}(v))$ ;
6   |  $v = j$ 'th out-neighbor of  $v$ ;
7 return  $\infty$ 

```

Q: What is the advantage of this algorithm over BFS?

A:

It can be shown that if G is an *undirected* graph with n vertices and m edges, then for an appropriate choice of $\ell = O(mn)$, with high probability $\text{RandomWalk}(G, s, \ell)$ will visit all vertices reachable from s . Thus, we obtain a *Monte Carlo* algorithm for UndirectedSTConnectivity.

Theorem 6.1. *UndirectedSTConnectivity can be solved by a Monte Carlo randomized algorithm with arbitrarily small error probability in time $O(mn)$ using only $O(1)$ words of memory in addition to the input.*