

1 Announcements

Recommended Reading: CLRS Sec 16.1–16.2

- Mid term Oct-17
- SRE 4 next Tuesday Oct-15.

2 Interval Scheduling

In Lecture 11, we modelled the conflicts arising in RAM allocation as a Coloring problem and discussed greedy algorithms for the problem. We found that the greedy - but careful - way of coloring via BFS worked well for 2-colorable graphs.

Another example where the greedy approach to algorithm design works well is Interval Scheduling. The following decision version of Interval Scheduling was introduced in Lecture 4.

Input	: A collection of intervals $[a_0, b_0], \dots, [a_{n-1}, b_{n-1}]$, where each $a_i, b_i \in \mathbb{Q}$ and $a_i \leq b_i$
Output	: YES if the intervals are disjoint (for all $i \neq j$, $[a_i, b_i] \cap [a_j, b_j] = \emptyset$) NO otherwise

Computational Problem IntervalScheduling-Decision

We saw that we could solve this problem in time $O(n \log n)$ by reduction to Sorting. However, if the answer is NO, we might be satisfied by trying to schedule *as many* intervals *as possible*:

Input	: A collection of intervals $[a_0, b_0], \dots, [a_{n-1}, b_{n-1}]$, where each $a_i, b_i \in \mathbb{Q}$ and $a_i \leq b_i$
Output	: A maximum-size subset $S \subseteq [n]$ such that $\forall i \neq j \in S$, $[a_i, b_i] \cap [a_j, b_j] = \emptyset$.

Computational Problem IntervalScheduling-Optimization

A greedy algorithm for IntervalScheduling-Optimization is as follows.

1 GreedyIntervalScheduling(x)	
Input	: A list x of n intervals $[a, b]$, with $a, b \in \mathbb{Q}$
Output	: A “large” subset of the input intervals that are disjoint from each other
2	Sort the intervals with ;
3	$S = \emptyset$;
4	foreach $i = 0$ to $n - 1$ do
5	;
6	return S

Theorem 2.1. GreedyIntervalScheduling(x) finds an optimal solution to IntervalScheduling-Optimization, and can be implemented in time $O(n \log n)$.

Proof.

□

3 Independent Set

In Lecture 11, we used graph theoretic modelling on RAM allocation to rephrase it as a Coloring problem. Graph theoretic modelling can also be applied to the IntervalScheduling-Optimization problem. For example, consider the following set of intervals:

This is the notion of an independent set, which is defined as follows for a general graph.

Definition 3.1. Let $G = (V, E)$ be a graph. An *independent set* in G is a subset $S \subseteq V$ such that there are no edges entirely in S . That is, $\{u, v\} \in E$ implies that $u \notin S$ or $v \notin S$.

Finding the largest set of conflict-free intervals thus amounts to finding the largest independent set in the corresponding conflict graph. For a general graphs, this problem is known as the Independent Set problem.

Input	: A graph $G = (V, E)$
Output	: An independent set $S \subseteq V$ in G of maximum size

Computational Problem Independent Set

Example:

Remarks:

- **Independent Set vs IntervalScheduling-Decision:**

- **Independent Set vs Coloring:**

There are no known efficient (polynomial-time) algorithms for the Independent Set problem; we will seek to understand why later in the course. However, a greedy algorithm along the lines of `GreedyIntervalScheduling` and `GreedyColoring` can be designed which outputs a somewhat large independent set.

```

1 GreedyIndSet( $G$ )
   Input      : A graph  $G = (V, E)$ 
   Output     : A “large” independent set in  $G$ 
2 Choose an ordering  $v_0, v_1, v_2, \dots, v_{n-1}$  of  $V$ ;
3  $S = \emptyset$ ;
4 foreach  $i = 0$  to  $n - 1$  do
5     |                                     ;
6     | return  $S$ 
```

How large is the independent set in the above algorithm? Similarly to coloring, we can only prove fairly weak bounds on the performance of the greedy algorithm in general:

Theorem 3.2. *For every graph G with n vertices and m edges, `GreedyIndSet`(G) can be implemented in time $O(n + m)$ and outputs an independent set of size at least $n/(\deg_{\max} + 1)$, where \deg_{\max} is the maximum vertex degree in G .*

Proof.

□

4 Matching

While Independent Set is a graph theoretic notion about “non-conflicting” vertices, matching is about “non-conflicting” edges.

Definition 4.1. For a graph $G = (V, E)$, a *matching* in G is a subset $M \subseteq E$ such that every vertex $v \in V$ is incident to at most one edge in M . Equivalently, no two edges in M share an endpoint.

If a vertex v is incident to an edge in M , we say v is *matched* by M ; otherwise we say it is *unmatched*.

The problem of finding the largest matching in a graph is called Maximum Matching.

Input	: A graph $G = (V, E)$
Output	: A matching $M \subseteq E$ in G of maximum size

Computational Problem Maximum Matching

An example of maximum matching is depicted below

We can use a greedy strategy to try finding a maximum matching. We can start with an edge and increase the size of the matching by adding edges (that do not share a vertex with edges that have already been added) till we can no longer add edges. An example is depicted below

Remarkably, sometimes it is possible to do more sophisticated operations than just adding an edge and still grow the matching. For this, we need the notions of alternating walk and augmenting path.

Definition 4.2. Let $G = (V, E)$ be a graph, and M be a matching in G . Then:

1. An *alternating walk* W in G with respect to M is
2. An *augmenting path* P in G with respect to M is

An example of alternating walk is depicted in Figure 1, which is also an augmenting path. We can further observe that the set of edges appearing in this augmenting path is precisely the *symmetric difference* between the matching on the left hand side of Figure 1 and the maximum matching in Figure ???. Thus, intuitively, the augmenting path tells us which edges to ‘switch’ to go from the matching of size 3 to the matching of size 4.




Figure 1: For the matching of size 3 on the left, an alternating walk on the right. This is also an augmenting path, since all the vertices are distinct and the first and the last vertices are not in the matching.

This suggests a natural algorithm for maximum matching: repeatedly try to find an augmenting path and use it to grow our matching. But we need to argue that augmenting paths always exist and we can find them efficiently. An important piece in this argument is the following theorem.

Theorem 4.3 (Berge’s Theorem). *Let $G = (V, E)$ be a graph, and $M \subseteq E$ be a matching. If (and only if) M is not a maximum-size matching, then G has an augmenting path with respect to M .*

In the next lecture, we will see how to turn the above idea into an algorithm for the case of *bipartite* graphs.

Definition 4.4. A graph $G = (V, E)$ is *bipartite* if it is 2-colorable. That is, there is a partition of vertices $V = V_0 \cup V_1$ (with $V_0 \cap V_1 = \emptyset$) such that all edges in E have one endpoint in V_0 and one endpoint in V_1 .

The example in Figure ??? is a bipartite graph.